

Netty - Best Practices a.k.a **Faster == Better**

@Twitter 2014 · San Francisco · 2014/07/28

Norman Maurer, Principal Software Engineer / Leading Netty efforts
@ Red Hat Inc.

  **Netty / All things NIO**

  **Author of Netty in Action**

  **@normanmaurer**

  **github.com/normanmaurer**

- i** ☐ **Pipelining**
- i** ☐ **Writing gracefully**
- i** ☐ **Buffers best-practises**
- i** ☐ **EventLoop**

No Pipelining Optimization

```
public class HttpHandler extends SimpleChannelInboundHandler<HttpRequest> {  
    @Override  
    public void channelRead(ChannelHandlerContext ctx, HttpRequest req) {  
        ChannelFuture future = ctx.writeAndFlush(createResponse(req)); ❶  
        if (!isKeepAlive(req)) {  
            future.addListener(ChannelFutureListener.CLOSE); ❷  
        }  
    }  
}
```

❶ Write to the Channel and flush out to the Socket.

❷ After written close Socket

Pipelining to safe syscalls!

```
public class HttpPipeliningHandler extends SimpleChannelInboundHandler<HttpRequest> {  
    @Override  
    public void channelRead(ChannelHandlerContext ctx, HttpRequest req) {  
        ChannelFuture future = ctx.writeAnd(createResponse(req)); ❶  
        if (!isKeepAlive(req)) {  
            future.addListener(ChannelFutureListener.CLOSE); ❷  
        }  
    }  
    @Override  
    public void channelReadComplete(ChannelHandlerContext ctx) {  
        ctx.flush(); ❸  
    }  
}
```

- ❶ **Write** to the `Channel` (**No syscall!**) but don't flush yet
- ❷ **Close** socket when done writing
- ❸ **Flush** out to the socket.

write(msg) , flush() and writeAndFlush(msg)

write(msg) \Rightarrow pass through pipeline

flush() \Rightarrow gathering write of previous written msgs

writeAndFlush() \Rightarrow short-cut for **write(msg)** and **flush()**


💡 Limit flushes as much as possible as syscalls are quite expensive.

💡 But also limit `write(...)` as much as possible as it need to traverse the whole pipeline.

May write too fast!

```
public class BlindlyWriteHandler extends ChannelInboundHandlerAdapter {  
    @Override  
    public void channelActive(ChannelHandlerContext ctx) throws Exception {  
        while(needsToWrite) { ❶  
            ctx.writeAndFlush(createMessage());  
        }  
    }  
}
```

❶ Writes till `needsToWrite` returns `false`.

 Risk of **OutOfMemoryError** if writing too fast and having slow receiver!

Correctly write with respect to slow receivers

```
public class GracefulWriteHandler extends ChannelInboundHandlerAdapter {  
    @Override  
    public void channelActive(ChannelHandlerContext ctx) {  
        writeIfPossible(ctx.channel());  
    }  
    @Override  
    public void channelWritabilityChanged(ChannelHandlerContext ctx) {  
        writeIfPossible(ctx.channel());  
    }  
  
    private void writeIfPossible(Channel channel) {  
        while(needsToWrite && channel.isWritable()) { ❶  
            channel.writeAndFlush(createMessage());  
        }  
    }  
}
```

❶ Make proper use of `Channel.isWritable()` to prevent `OutOfMemoryError`

Configure high and low write watermarks

💡 Set sane **WRITE_BUFFER_HIGH_WATER_MARK** and **WRITE_BUFFER_LOW_WATER_MARK**

Server

```
ServerBootstrap bootstrap = new ServerBootstrap();  
bootstrap.childOption(ChannelOption.WRITE_BUFFER_HIGH_WATER_MARK, 32 * 1024);  
bootstrap.childOption(ChannelOption.WRITE_BUFFER_LOW_WATER_MARK, 8 * 1024);
```

Client

```
Bootstrap bootstrap = new Bootstrap();  
bootstrap.option(ChannelOption.WRITE_BUFFER_HIGH_WATER_MARK, 32 * 1024);  
bootstrap.option(ChannelOption.WRITE_BUFFER_LOW_WATER_MARK, 8 * 1024);
```

Issues with using non pooled buffers

 Use unpooled buffers with **caution!**

 **Allocation / Deallocation is slow**

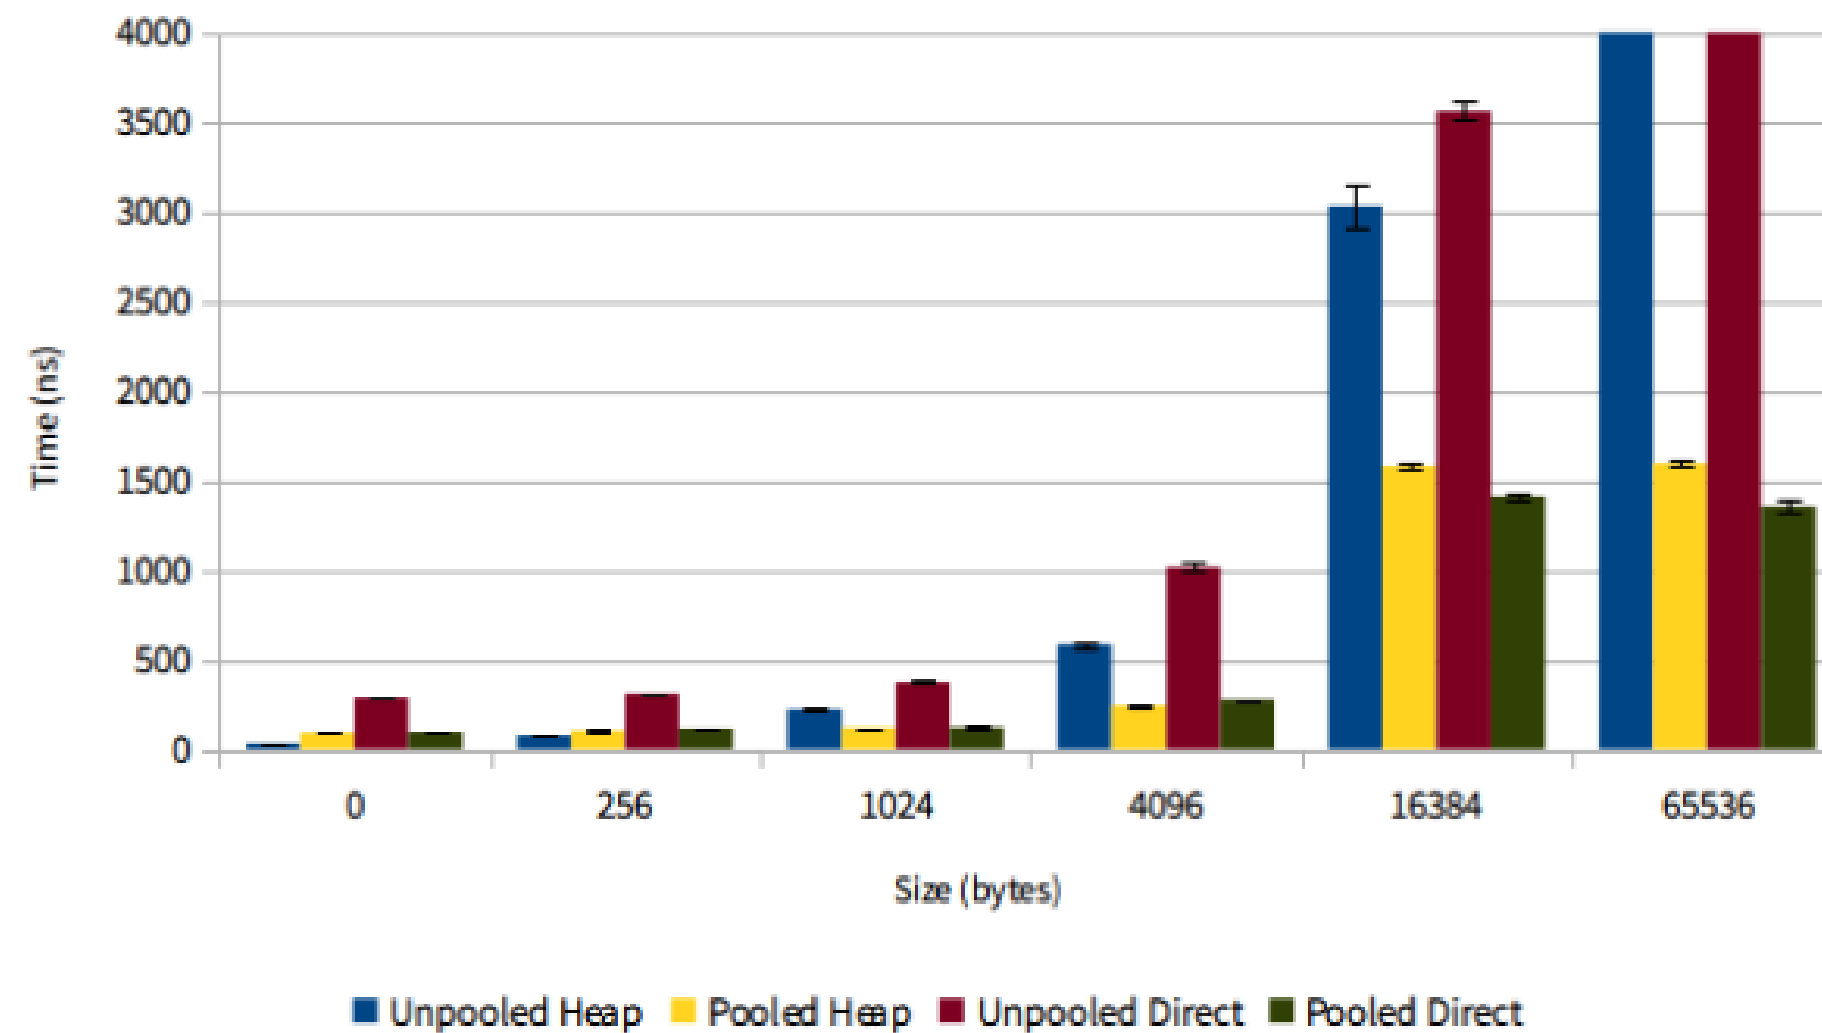
 **Free up direct buffers == PITA!**

 **Use pooled buffers!**

```
Bootstrap bootstrap = new Bootstrap();  
bootstrap.option(ChannelOption.ALLOCATOR, PooledByteBufAllocator.DEFAULT);  
ServerBootstrap bootstrap = new ServerBootstrap();  
bootstrap.childOption(ChannelOption.ALLOCATOR, PooledByteBufAllocator.DEFAULT);
```


Use Pooling of buffers to reduce allocation / deallocation time!

💡 Pooling pays off for direct and heap buffers!



<https://blog.twitter.com/2013/netty-4-at-twitter-reduced-gc-overhead>

Always use direct ByteBuffer when writing to SocketChannel

 OpenJDK and Oracle JDK copy otherwise to direct buffer by itself!

Only use heap buffers if need to operate on byte[] in `ChannelOutboundHandler`! By default direct ByteBuffer will be returned by `ByteBufferAllocator.buffer(...)`.

Take this as rule of thumb

Find pattern in ByteBuf

SlowSearch :(

```
ByteBuf buf = ...;
int index = -1;
for (int i = buf.readerIndex(); index == -1 && i < buf.writerIndex(); i++) {
    if (buf.getBytes(i) == '\n') {
        index = i;
    }
}
```

FastSearch :)

```
ByteBuf buf = ...;
int index = buf.forEachByte(new ByteBufProcessor() {
    @Override
    public boolean process(byte value) {
        return value != '\n';
    }
});
```

Messages with Payload? Yes please...

ByteBuffer payload \Rightarrow extend **DefaultByteBufferHolder**

i **reference-counting for free**

i **release resources out-of-the-box**







<http://www.flickr.com/photos/za3tooor/65911648/>

💡 Use zero-memory-copy for efficient transfer of raw file content

```
Channel channel = ...;  
FileChannel fc = ...;  
channel.writeAndFlush(new DefaultFileRegion(fc, 0, fileLength));
```

🔥 This only works if you don't need to modify the data on the fly. If so use `ChunkedWriteHandler` and `NioChunkedFile`.

Never block the EventLoop!

-  `Thread.sleep()`
-  `CountDownLatch.await()` or any other blocking operation from `java.util.concurrent`
-  Long-lived computationally intensive operations
-  Blocking operations that might take a while (e.g. DB query)



Re-use EventLoopGroup if you can!

```
Bootstrap bootstrap = new Bootstrap().group(new NioEventLoopGroup());  
Bootstrap bootstrap2 = new Bootstrap().group(new NioEventLoopGroup());
```

Share EventLoopGroup between different Bootstraps

```
EventLoopGroup group = new NioEventLoopGroup();  
Bootstrap bootstrap = new Bootstrap().group(group);  
Bootstrap bootstrap2 = new Bootstrap().group(group);
```

💡 **Sharing** the same **EventLoopGroup** allows to keep the resource usage (like Thread-usage) to a minimum.

Proxy like application with context-switching issue

```
public class ProxyHandler extends ChannelInboundHandlerAdapter {  
    @Override  
    public void channelActive(ChannelHandlerContext ctx) { ❶  
        final Channel inboundChannel = ctx.channel();  
        Bootstrap b = new Bootstrap();  
        b.group(new NioEventLoopGroup()); ❷  
        ...  
        ChannelFuture f = b.connect(remoteHost, remotePort);  
        ...  
    }  
}
```

❶ Called once a new connection was accepted

❷ Use a new `EventLoopGroup` instance to handle the connection to the remote peer

 Don't do this! This will tie up more resources than needed and introduce extra context-switching overhead.

Proxy like application which reduce context-switching to minimum

```
public class ProxyHandler extends ChannelInboundHandlerAdapter {  
    @Override  
    public void channelActive(ChannelHandlerContext ctx) { ❶  
        final Channel inboundChannel = ctx.channel();  
        Bootstrap b = new Bootstrap();  
        b.group(inboundChannel.eventLoop()); ❷  
        ...  
        ChannelFuture f = b.connect(remoteHost, remotePort);  
        ...  
    }  
}
```

❶ Called once a new connection was accepted

❷ Share the same `EventLoop` between both Channels. This means all IO for both connected Channels are handled by the same Thread.

💡 Always **share** EventLoop in those Applications

Operations from inside ChannelHandler

```
public class YourHandler extends ChannelInboundHandlerAdapter {  
    @Override  
    public void channelActive(ChannelHandlerContext ctx) {  
        // BAD (most of the times)  
        ctx.channel().writeAndFlush(msg); ❶  
  
        // GOOD  
        ctx.writeAndFlush(msg); ❷  
    }  
}
```

- ❶ `Channel.*` methods \Rightarrow the operation will start at the tail of the `ChannelPipeline`
- ❷ `ChannelHandlerContext.*` methods \Rightarrow the operation will start from this `ChannelHandler` to flow through the `ChannelPipeline`.

💡 Use the shortest **path** as possible to get the maximal performance.

Share ChannelHandlers if stateless

```
@ChannelHandler.Shareable ❶
public class StatelessHandler extends ChannelInboundHandlerAdapter {
    @Override
    public void channelActive(ChannelHandlerContext ctx) {
        logger.debug("Now client from " + ctx.channel().remoteAddress().toString());
    }
}

public class MyInitializer extends ChannelInitializer<Channel> {
    private static final ChannelHandler INSTANCE = new StatelessHandler();
    @Override
    public void initChannel(Channel ch) {
        ch.pipeline().addLast(INSTANCE);
    }
}
```

- ❶ Annotate `ChannelHandler` that are stateless with `@ChannelHandler.Shareable` and use the same instance across Channels to reduce GC.

Remove ChannelHandler once not needed anymore

```
public class OneTimeHandler extends ChannelInboundHandlerAdapter {  
    @Override  
    public void channelActive(ChannelHandlerContext ctx) {  
        doOneTimeAction();  
        ctx.channel().pipeline().remove(this); ❶  
    }  
}
```

❶ Remove `ChannelHandler` once not needed anymore.

💡 This keeps the `ChannelPipeline` as **short** as possible and so **eliminate overhead** of traversing as much as possible.

Use proper buffer type in MessageToByteEncoder

```
public class EncodeActsOnByteArray extends MessageToByteEncoder<YourMessage> {  
    public EncodeActsOnByteArray() { super(false); } ❶  
    @Override  
    public encode(ChannelHandlerContext ctx, YourMessage msg, ByteBuf out) {  
        byte[] array = out.array(); ❷  
        int offset = out.arrayOffset() + out.writerIndex();  
        out.writeIndex(out.writerIndex() + encode(msg, array, offset)); ❸  
    }  
    private int encode(YourMessage msg, byte[] array, int offset, int len) { ... }  
}
```

- ❶ Ensure **heap buffers** are used when pass into `encode(...)` method. This way you can access the backing array directly
- ❷ Access the **backing array** and also calculate offset
- ❸ Update **writerIndex** to reflect written bytes

💡 This saves extra byte copies.

To auto-read or not to auto-read

By default Netty will keep on reading data from the `Channel` once something is ready.

To auto-read or not to auto-read

By default Netty will keep on reading data from the **Channel** once something is ready.

Need more fine grained control ?


```
channel.config().setAutoRead(false); ❶  
channel.read(); ❷  
channel.config().setAutoRead(true); ❸
```

- ❶ Disable auto read == no more data will be read automatically from this **Channel**.
- ❷ Tell the **Channel** to do one read operation once new data is ready
- ❸ Enable again auto read == Netty will automatically read again

💡 This can also be quite useful when writing proxy like applications!

Native stuff in Netty 4

- 💡 OpenSSL based SslEngine to reduce memory usage and latency.
- 💡 Native transport for Linux using Epoll ET for more performance and less CPU usage.
- 💡 Native transport also supports SO_REUSEPORT and TCP_CORK :)



Switching to native transport is easy

Using NIO transport


```
Bootstrap bootstrap = new Bootstrap().group(new NioEventLoopGroup());  
bootstrap.channel(NioSocketChannel.class);
```

Using native transport

```
Bootstrap bootstrap = new Bootstrap().group(new EpollEventLoopGroup());  
bootstrap.channel(EpollSocketChannel.class);
```

Want to know more?

💡 Buy my book [Netty in Action](#) and make me **RICH**.



<http://www.manning.com/maurer>

\$ KA-CHING \$

References

- i Netty - <http://netty.io>
 - i Slides generated with Ascidoctor and DZSlides backend
 - i Original slide template - Dan Allen & Sarah White
 - i All pictures licensed with `Creative Commons Attribution` or `Creative Commons Attribution-Share Alike`
-

Norman Maurer

  @normanmaurer