

# Why Netty ?

---

@Netflix 2014 · Las Gatos · 2014/07/29

**Norman Maurer, Principal Software Engineer / Leading Netty efforts  
@ Red Hat Inc.**

-  **Netty / All things NIO**
-  **Author of Netty in Action**
-  **@normanmaurer**
-  **github.com/normanmaurer**

## Netty - The non-blocking network framework

“Because writing fast, non-blocking  
network code is non-trivial.

– Why Netty ?

# When I say fast, I mean fast

Best (bar chart)	Data table	Latency	Framework overhead	Best plaintext responses per second, i7-2600K hardware (67 tests)							
Framework	Best performance (higher is better)				Cls	Lng	Plt	FE	Aos	IA	Errors
grizzly	631,280	I	100.0%	Mcr	Jav	Svt	Grz	Lin	Rea	0	
netty	624,786	I	99.0%	Plt	Jav	Nty	Non	Lin	Rea	0	
undertow	612,146	I	97.0%	Plt	Jav	Utw	Non	Lin	Rea	197	
undertow edge	610,948	I	96.8%	Plt	Jav	Und	Non	Lin	Rea	50	

<http://www.techempower.com/benchmarks/#section=data-r9&hw=i7&test=plaintext>

## Response

HTTP/1.1 200 OK

Content-Length: 15

Content-Type: text/plain; charset=UTF-8

Server: Example

Date: Wed, 17 Apr 2013 12:00:00 GMT

Hello, World!

## Even more speed in upcoming release

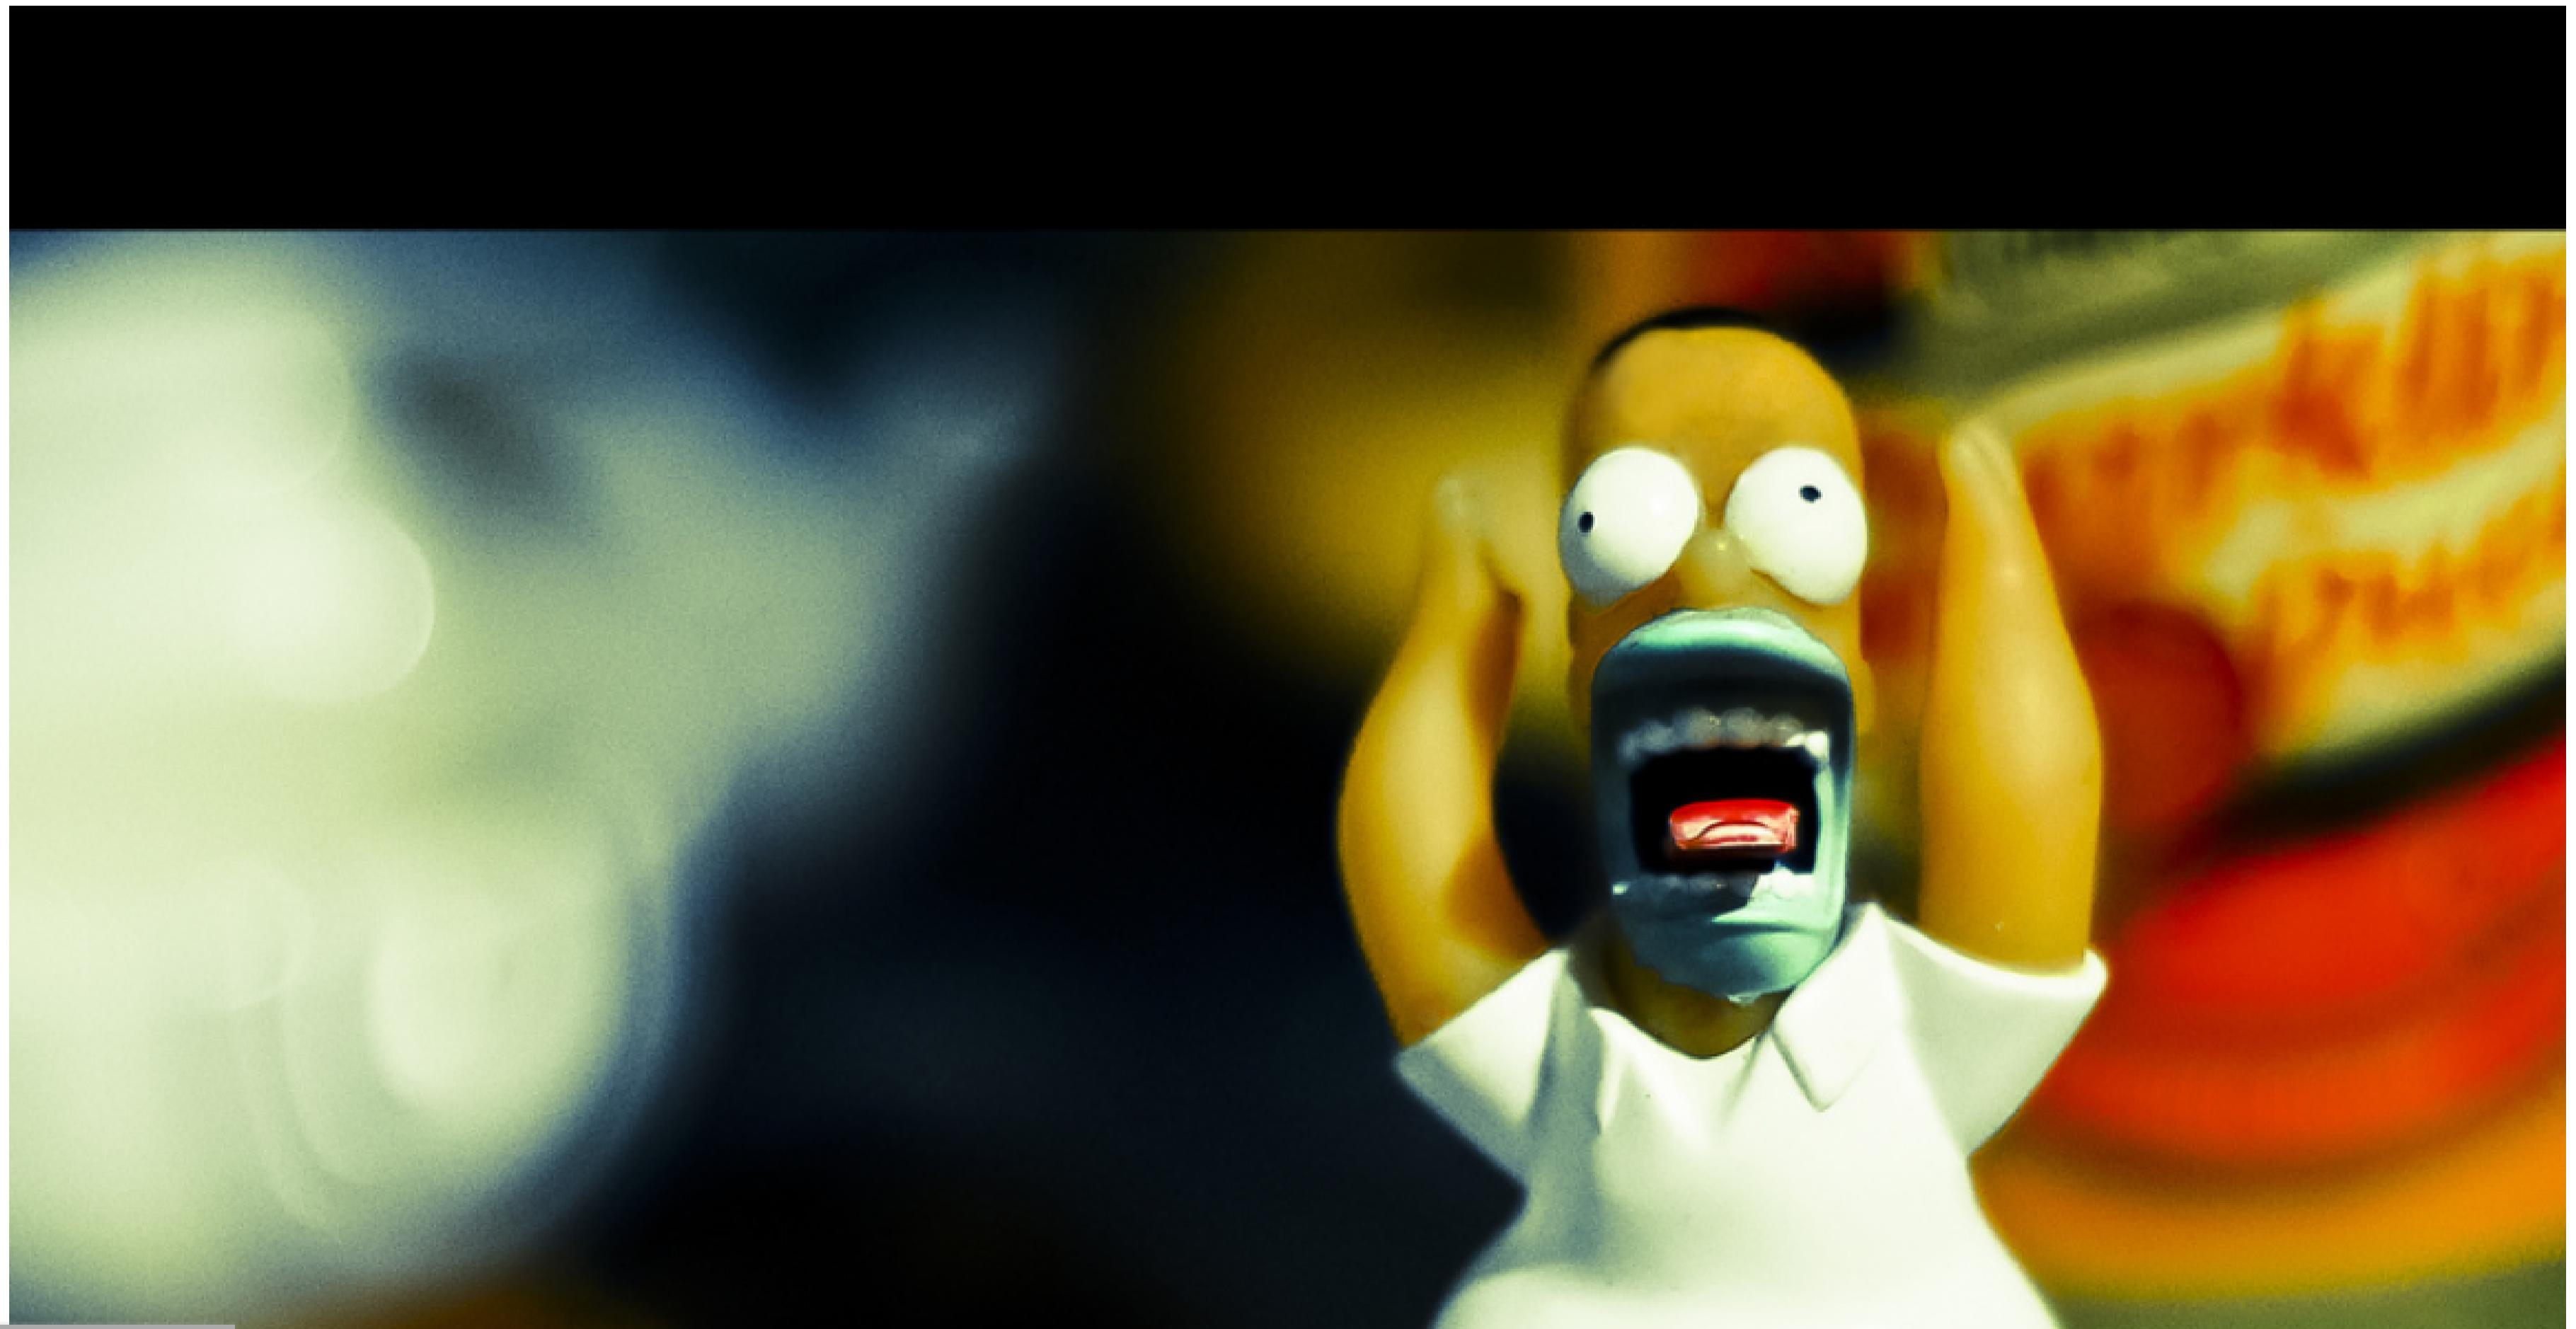
**i** 256 concurrent connections

**i** 256 requests pipelined

**24 cores**

```
[nmaurer@xxx]~% wrk/wrk -H 'Host: localhost' -H 'Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8' -H 'Connection: keep-alive' -d 120 -c 256 -t 16 --pipeline 256 http://xxx:8080/plaintext
Running 2m test @ http://xxx:8080/plaintext
  16 threads and 256 connections
  Thread Stats      Avg      Stdev      Max  +/- Stdev
    Latency    18.77ms   16.68ms  452.00ms   92.67%
    Req/Sec   225.82k   41.95k  376.21k   67.34%
  429966998 requests in 2.00m, 57.66GB read
  Requests/sec: 3583411.40
  Transfer/sec: 492.11MB
```

**But.... WHY?**



## Fully asynchronous

- ➊ Asynchronous from the ground up
- ➋ Using java.nio or native method calls for non-blocking io
- ➌ Futures and callbacks provided for easy composing

“

*Don't call us, we'll call you.*

– Hollywood principle

## Hide complexity but not flexibility

- i** Hides all the complexity involved when you use `java.nio` or `java.nio2`
  - i** Still powers you with a lot of flexibility
  - i** Unified API for every transport
  - i** Allows easy testing of your custom code.
-

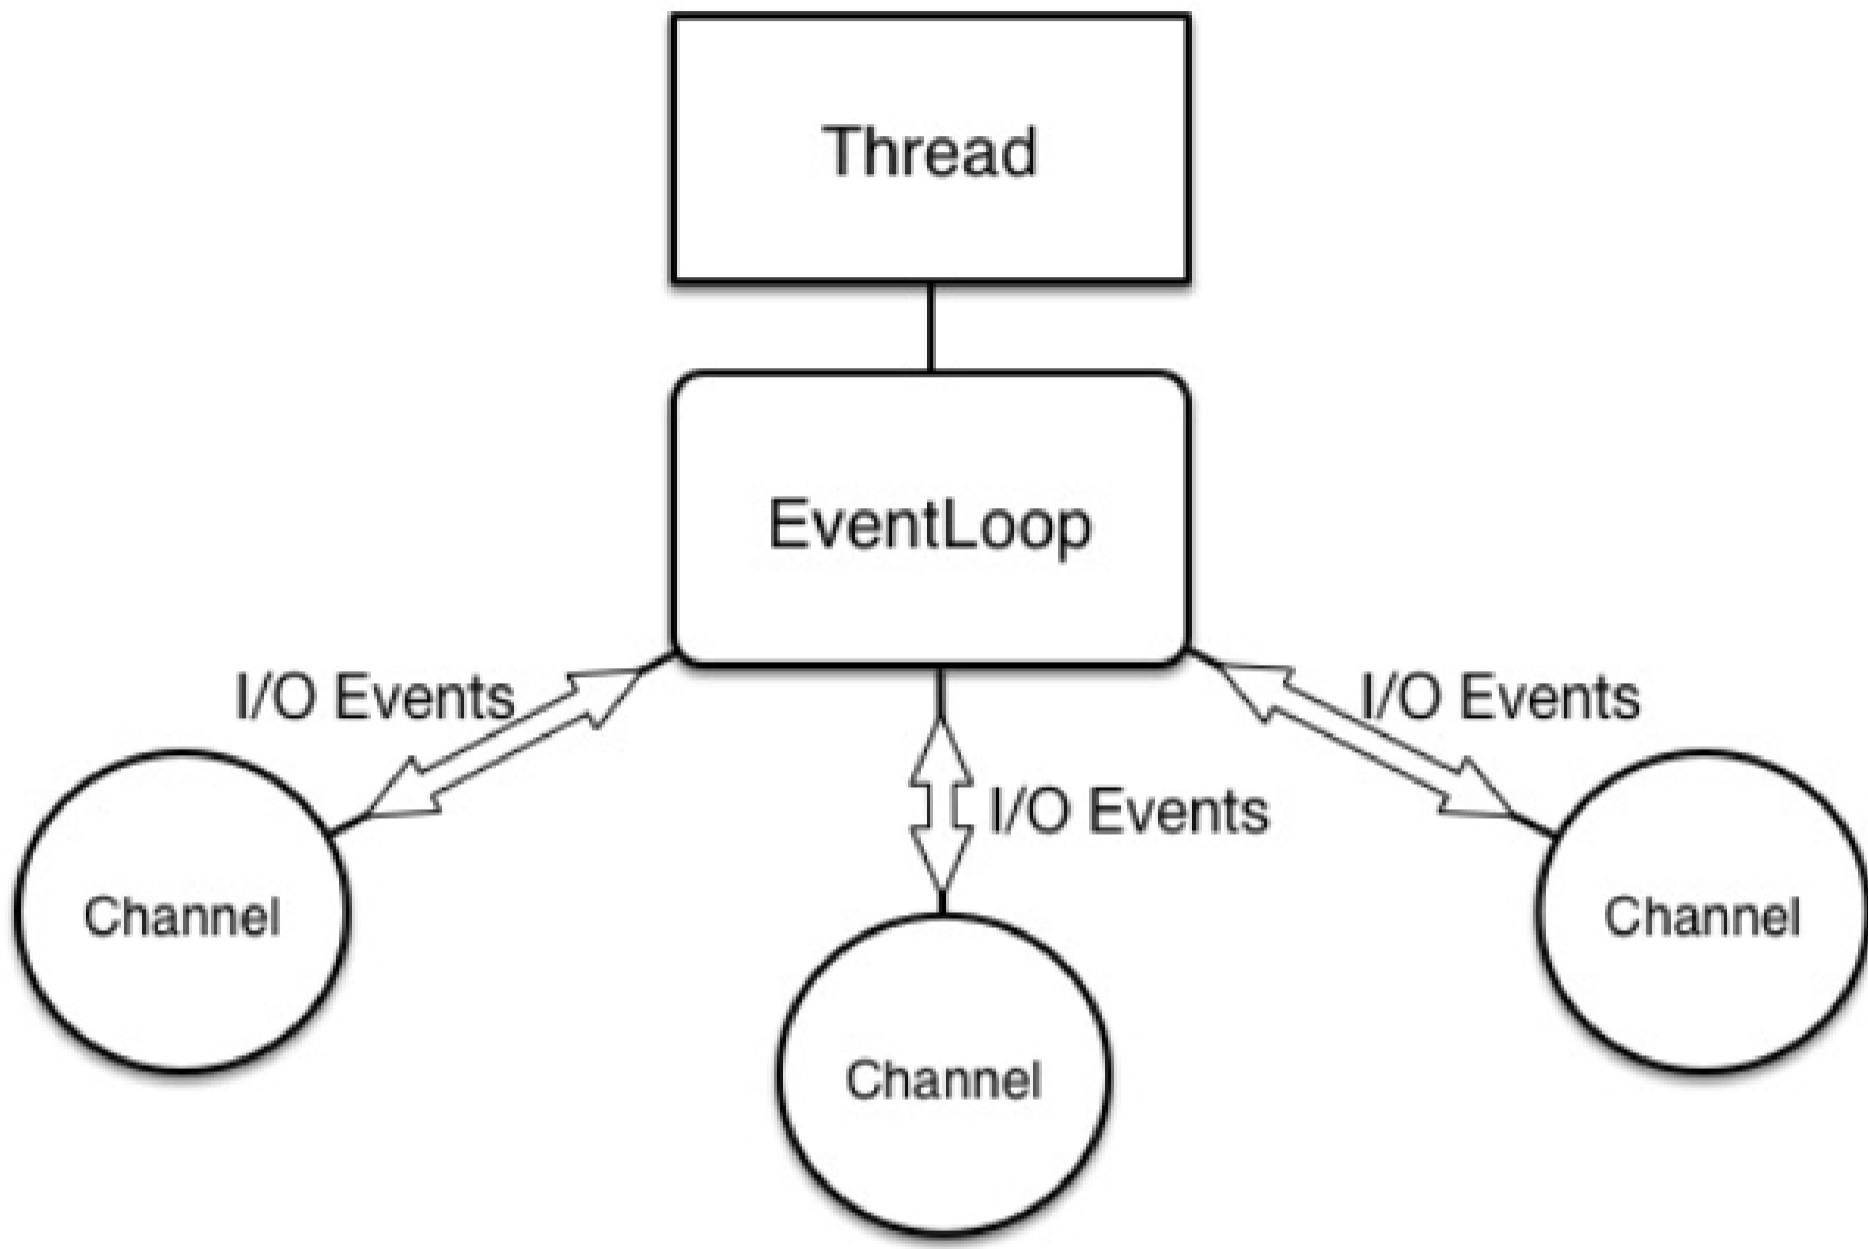
## Protocol agnostic → Not just another HTTP server

- Supports **TCP**, **UDP**, **UDT**, **SCTP** out of the box
- Contains codecs for different protocols on top of these

## Supported Codecs

- i** HTTP, WebSockets ( + compression) , SPDY, HTTP 2
- i** SSL/TLS, Zlib, Deflate
- i** Protobufs, JBoss Marshalling, Java Serialization
- i** DNS
- i** Memcached, Stomp, Proxy, MQTT
- 💡** ...add your preferred protocol here...

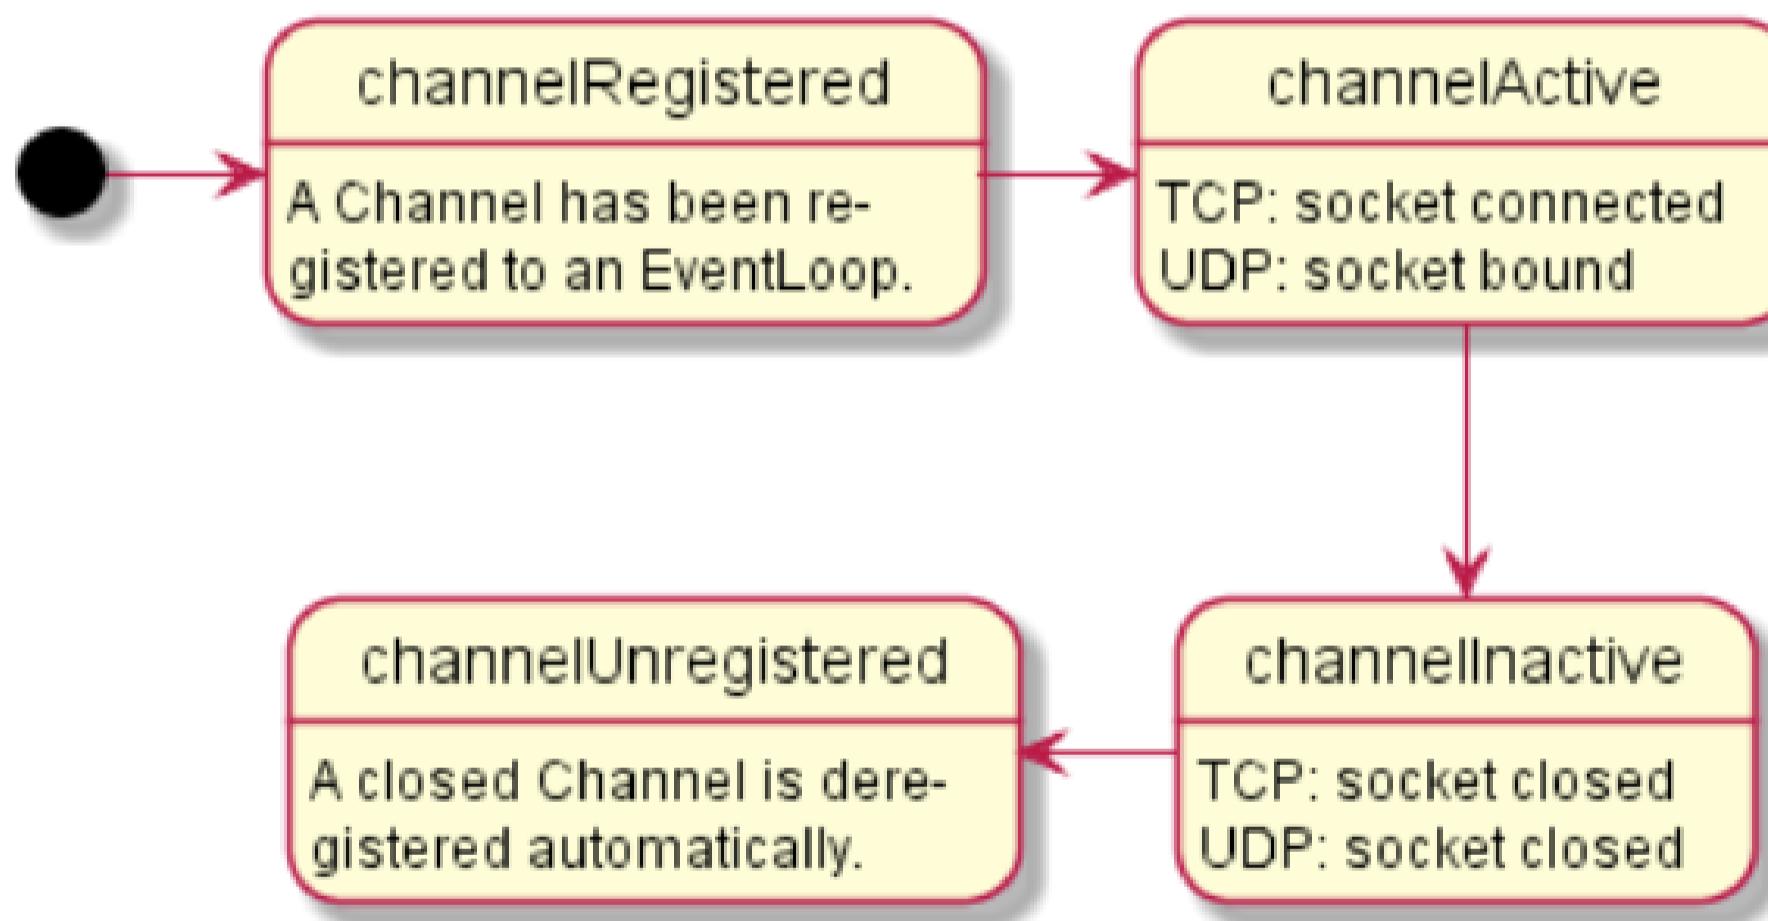
## Thread-Model - Easy but powerful



- 💡 Having inbound and outbound events handled by the same Thread simplifies concurrency handling a lot!

## Simple state model

- Allows to react on each state change by intercept the states via `ChannelHandler`.
- Allows flexible handling depending on the needs.



## ⓘ Interceptor pattern

- ⓘ Allows to add building-blocks (`ChannelHandler`) on-the-fly that transform data or react on events

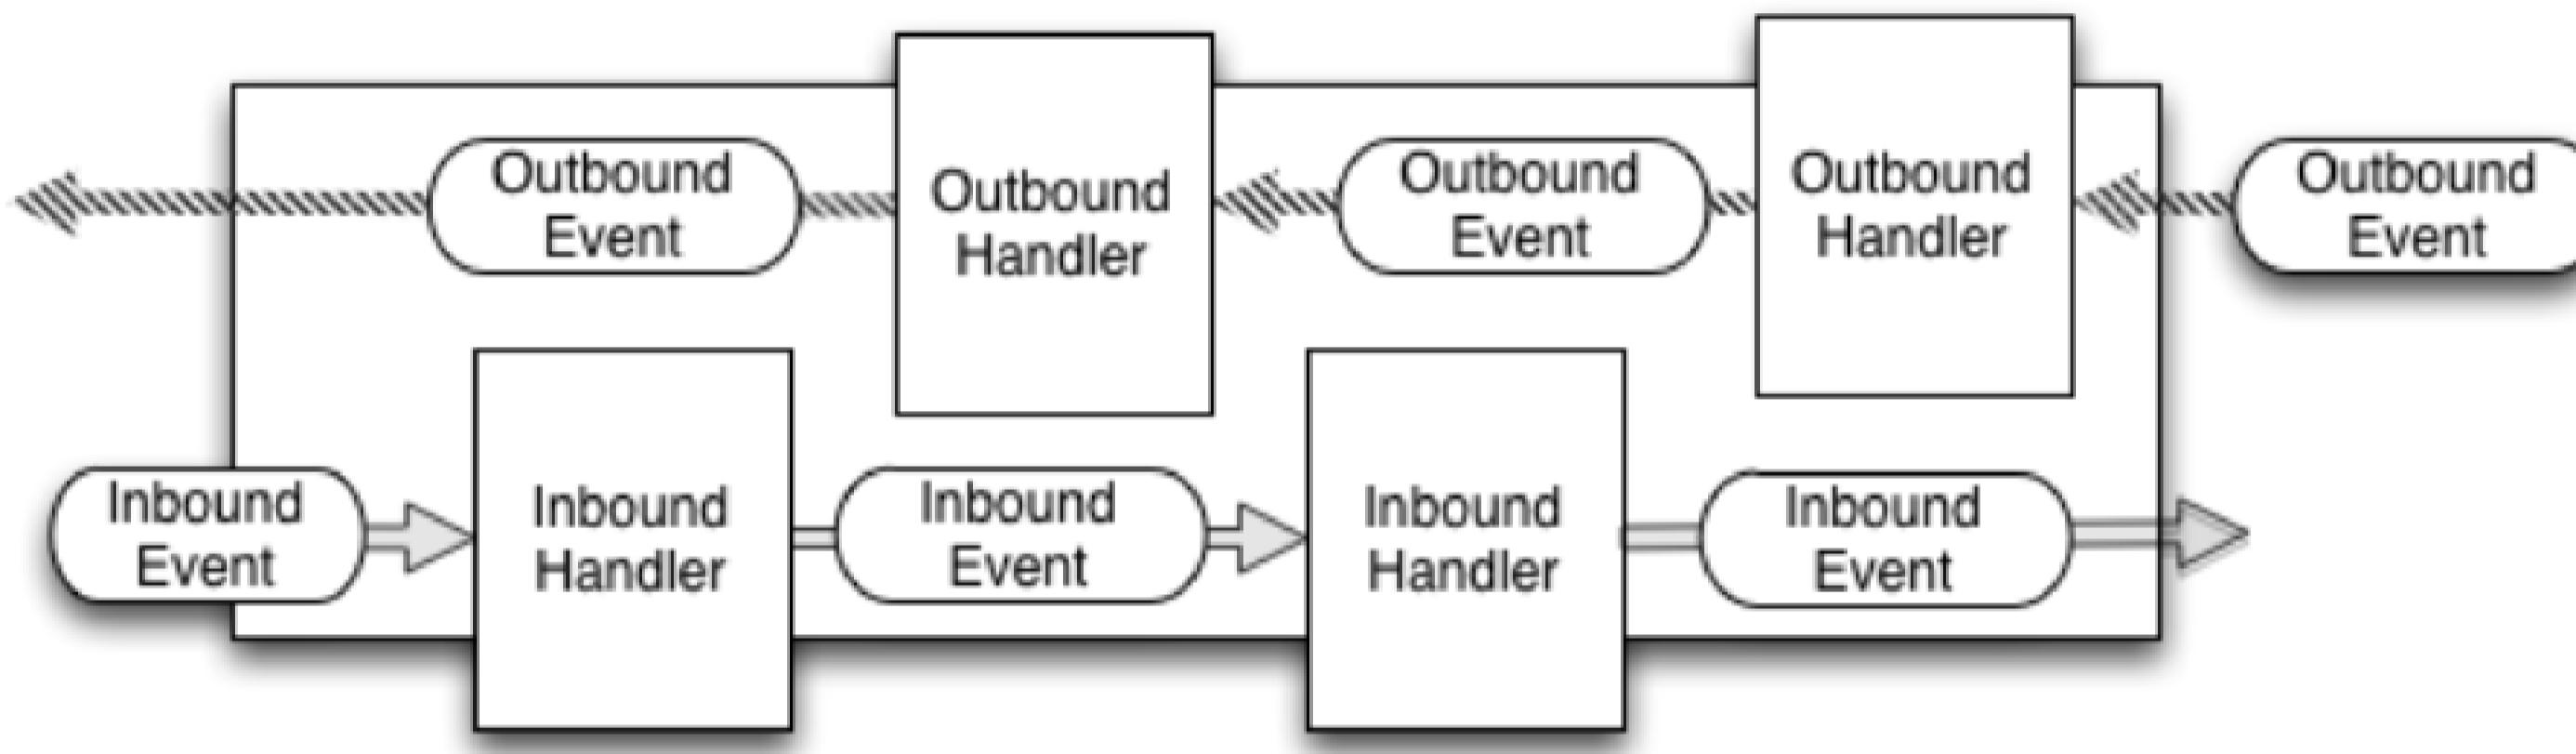
Kind of a unix-pipe-like thing...

```
$ echo "Netty is shit...." | sed -e 's/is /is the /' | cat ①  
Netty is the shit....
```

- ① Think of the whole line to be the `ChannelPipeline` and `echo`, `sed` and `cat` the `ChannelHandler`s that allow to transform data.

## ChannelPipeline - How does it work

- i Inbound and outbound events flow through the `ChannelHandler`'s in the `ChannelPipeline` and so allow to react one these events.



# ChannelPipeline - Compose processing logic

 Compose complex processing logic via multiple **ChannelHandler**.

```
public class MyChannelInitializer extends ChannelInitializer<Channel> {  
    @Override  
    public void initChannel(Channel ch) {  
        ChannelPipeline p = ch.pipeline();  
        p.addLast(new SslHandler(...)); ①  
        p.addLast(new HttpServerCodec(...)); ②  
        p.addLast(new YourRequestHandler()); ③  
    }  
}
```

- ① Encrypt traffic
- ② Support HTTP
- ③ Your handler that receive the HTTP requests.

# ChannelHandler - React on received data

```
@Sharable
public class EchoHandler extends ChannelInboundHandlerAdapter {
    @Override
    public void channelRead(ChannelHandlerContext ctx, Object msg) { ①
        ctx.writeAndFlush(msg);
    }

    @Override
    public void exceptionCaught(ChannelHandlerContext ctx, Throwable cause) { ②
        cause.printStackTrace();
        ctx.close();
    }
}
```

① Intercept received message and write it back to the remote peer

② React on **Throwable** and close the connection

## **Decoder / Encoder - Transform data via ChannelHandler**

- i** Different abstract base classes for Decoder and Encoder
- i** Handles buffering for you if needed (remember everything is non-blocking!)

# Decoder / Encoder - Transform data via ChannelHandler

## Transform received ByteBuf to String

```
public class StringDecoder extends MessageToMessageDecoder<ByteBuf> {  
    @Override  
    protected void decode(ChannelHandlerContext ctx, ByteBuf msg, List<Object> out) {  
        out.add(msg.toString(charset));  
    }  
}
```

## Transform to be send String to ByteBuf

```
public class StringEncoder extends MessageToMessageEncoder<String> {  
    @Override  
    protected void encode(ChannelHandlerContext ctx, String msg, List<Object> out) {  
        if (msg.length() == 0) return;  
        out.add(ByteBufUtil.encodeString(ctx.alloc(), CharBuffer.wrap(msg), charset));  
    }  
}
```

## Adding other processing logic?

- 💡 Adding more processing logic is often a matter of adding just-another `ChannelHandler` to the `ChannelPipeline`.



# Writing protocol multiplexers is a no-brainer!

```
public class PortUnificationServerHandler extends ByteToMessageDecoder {  
    @Override  
    protected void decode(ChannelHandlerContext ctx, ByteBuf in, List<Object> out) {  
        if (in.readableBytes() < 5) return ①;  
        if (isSsl(in)) { ②  
            ctx.pipeline().addLast("ssl", sslCtx.newHandler(ctx.alloc()));  
            ctx.pipeline().remove(this);  
        } else if (isGzip(in)) { ...  
        } else { ③  
            ctx.close();  
        }  
    }  
}
```

- ① Will use the first five bytes to detect a protocol.
- ② Check if SSL is used and if so add `SslHandler` to the `ChannelPipeline`.
- ③ Unknown protocol, just close the connection

## Flexible write behaviour

- i** `Channel.write(...)` ⇒ write through the `ChannelPipeline` but NOT trigger syscall like `write` or `writev`.
  
- i** `channel.flush()` ⇒ writes all pending data to the socket.
  
- 💡** Gives more flexibility for when things are written and also allows efficient pipelining.

# Easy HTTP Pipelining to safe syscalls by using write(...) and flush()!

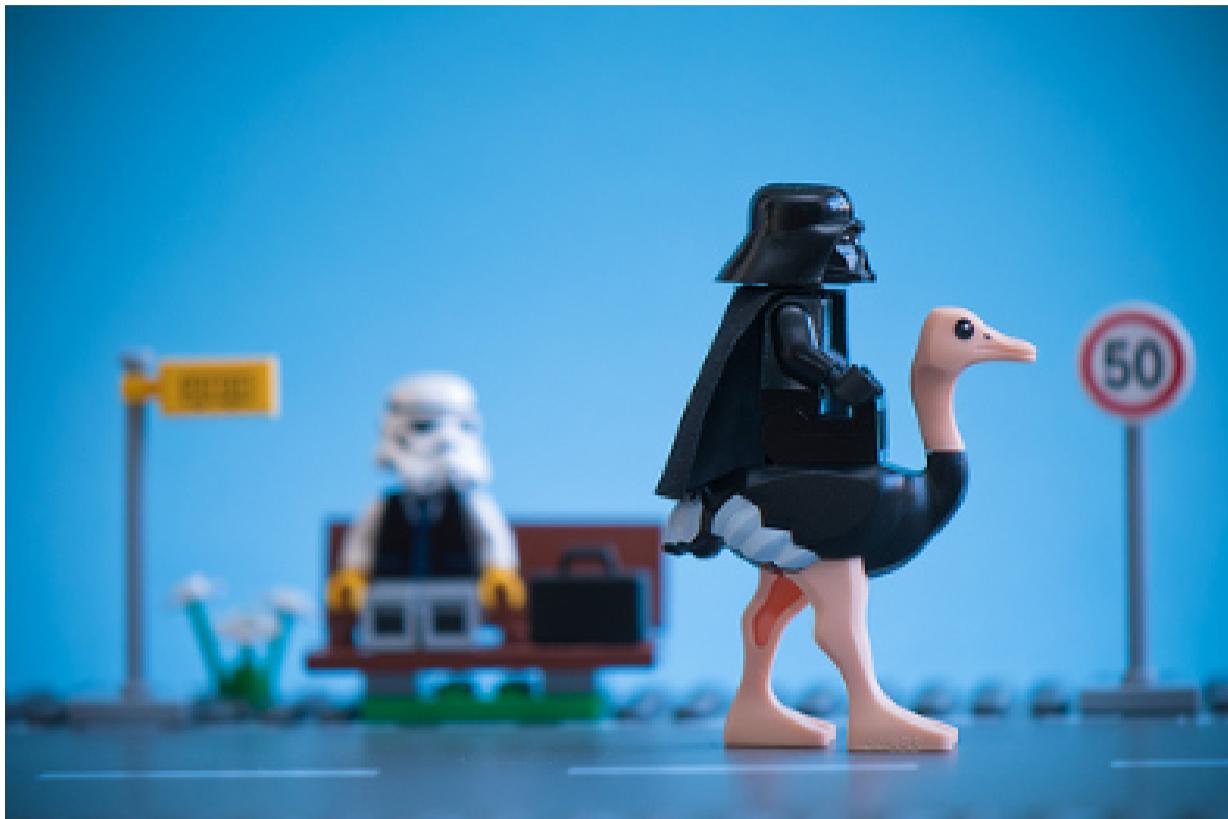
```
public class HttpPipeliningHandler extends SimpleChannelInboundHandler<HttpRequest> {  
    @Override  
    public void channelRead(ChannelHandlerContext ctx, HttpRequest req) {  
        ChannelFuture future = ctx.writeAndFlush(createResponse(req)); ①  
        if (!HttpHeaders.isKeepAlive(req)) {  
            future.addListener(ChannelFutureListener.CLOSE); ②  
        }  
    }  
    @Override  
    public void channelReadComplete(ChannelHandlerContext ctx) {  
        ctx.flush(); ③  
    }  
}
```

① Write to the `Channel` (**No syscall!**) but not flush yet

② After written `close` socket

③ Flush out to the socket once everything was ready from the 'Channel'

## Allow to detect slow remote peers



<https://www.flickr.com/photos/kwl/4514986410>

```
public class StateHandler extends ChannelInboundHandlerAdapter {  
    @Override  
    public void channelWritabilityChanged(ChannelHandlerContext ctx) { } ①  
}
```

① Is triggered once `Channel.isWritable()` changes.

## Execute ChannelHandler outside of EventLoop

- ⚠ Don't block as these will effect all `Channel`'s that are served by the same `Thread`.
- 💡 ChannelPipeline allows to add `ChannelHandler` that are executed on different Thread to free up IO Thread (`EventLoop`).

```
Channel ch = ...;  
ChannelPipeline p = ch.pipeline();  
EventExecutor e1 = new DefaultEventExecutor(16);
```

```
p.addLast(new MyProtocolCodec()); ①  
p.addLast(e1, new MyDatabaseAccessingHandler()); ②
```

① Executed in `EventLoop` (and so the `Thread` bound to it)

② Executed in one of the `EventExecutors` of e1

## Built with GC pressure in mind

- ❶ Use pooling / `ThreadLocal`s to prevent GC pressure
- ❶ Prefer direct method invocation over fire event object



[http://25.media.tumblr.com/tumblr\\_me2eq0PnBx1rtu0cp01\\_1280.jpg](http://25.media.tumblr.com/tumblr_me2eq0PnBx1rtu0cp01_1280.jpg)

- ❶ Reduces GC-Pressure a lot!

## ByteBuf - ByteBuffer on steroids!

- Separate index for reader/writer
  - Direct, Heap and Composite
  - Resizable with max capacity
  - Reference counting / pooling
- 💡 uses `sun.misc.Unsafe` for maximal performance

---

```
ByteBuf buf = ...;
buf.writeInt(1).writeBytes(data).writeBoolean(true)...
```

## Allow parsing without expensive range-checks

### SlowSearch for ByteBuf : (

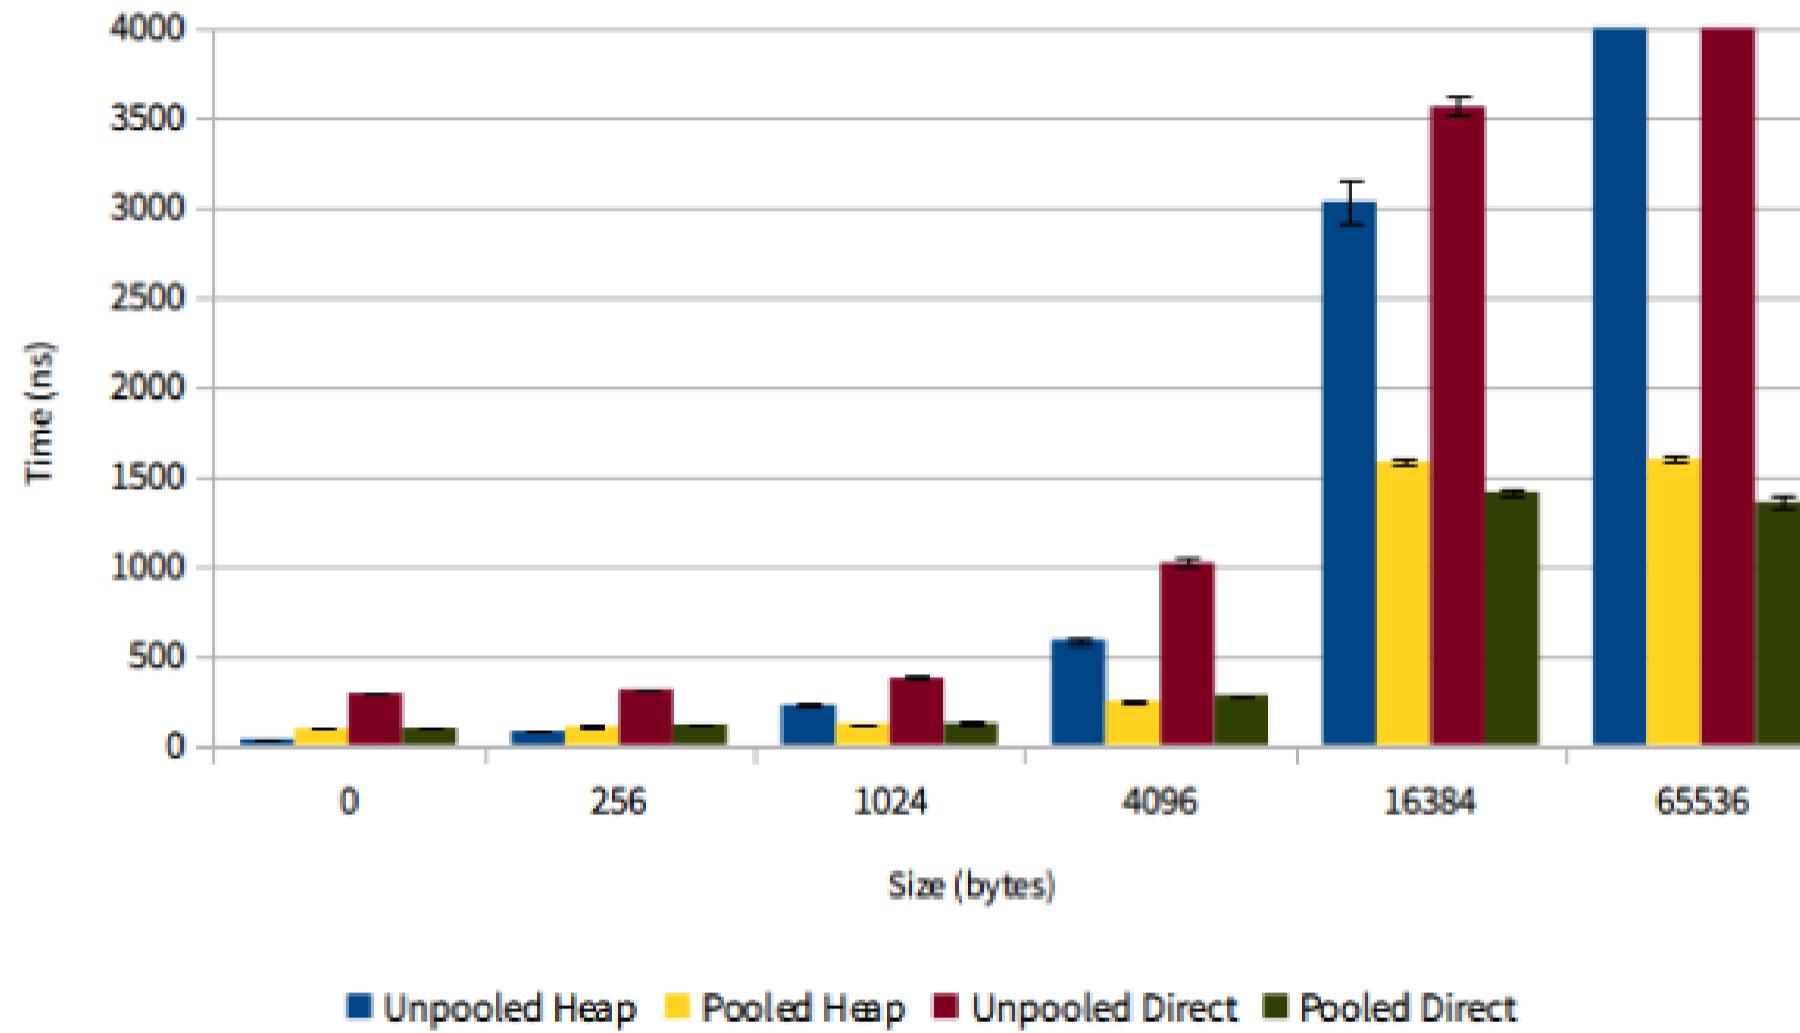
```
int index = -1;
for (int i = buf.readerIndex(); index == -1 && i < buf.writerIndex(); i++) {
    if (buf.getByte(i) == '\n') {
        index = i;
    }
}
```

### FastSearch for ByteBuf : )

```
int index = buf.forEachByte(new ByteBufProcessor() {
    @Override
    public boolean process(byte value) {
        return value != '\n';
    }
});
```

# Use Pooling of buffers to reduce allocation / deallocation time!

💡 Pooling pays off for direct and heap buffers!



<https://blog.twitter.com/2013/netty-4-at-twitter-reduced-gc-overhead>

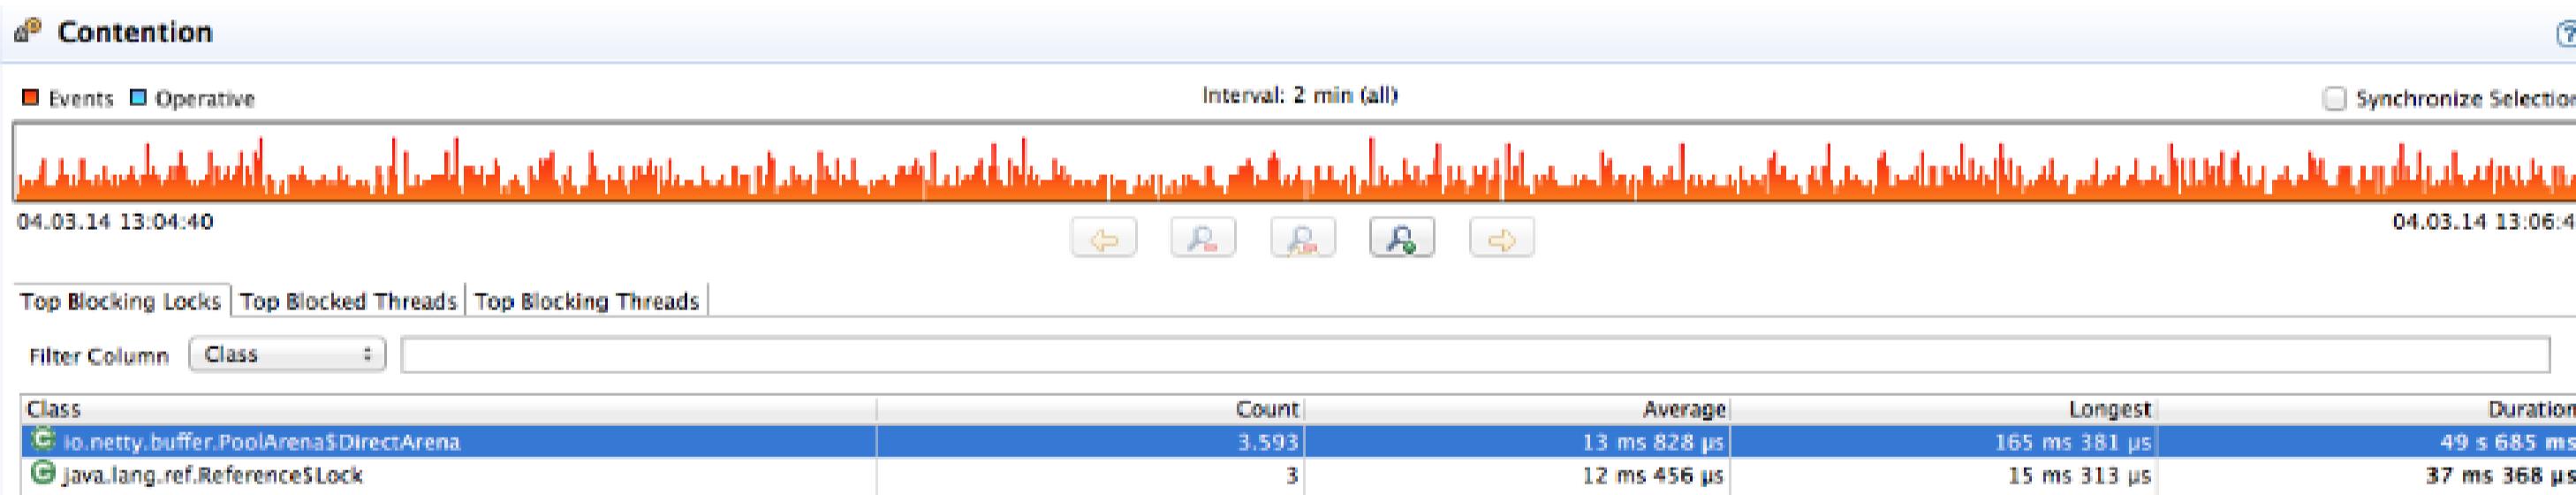
## PooledByteBufAllocator - Algorithms / datastructures

- Algorithm is a hybrid of jemalloc and buddy-allocation
- ThreadLocal caches for lock-free allocation
- Synchronization per area that holds different chunks of memory, when not be able to serve via cache

# PooledByteBufAllocator - without caches

## Before caches were added

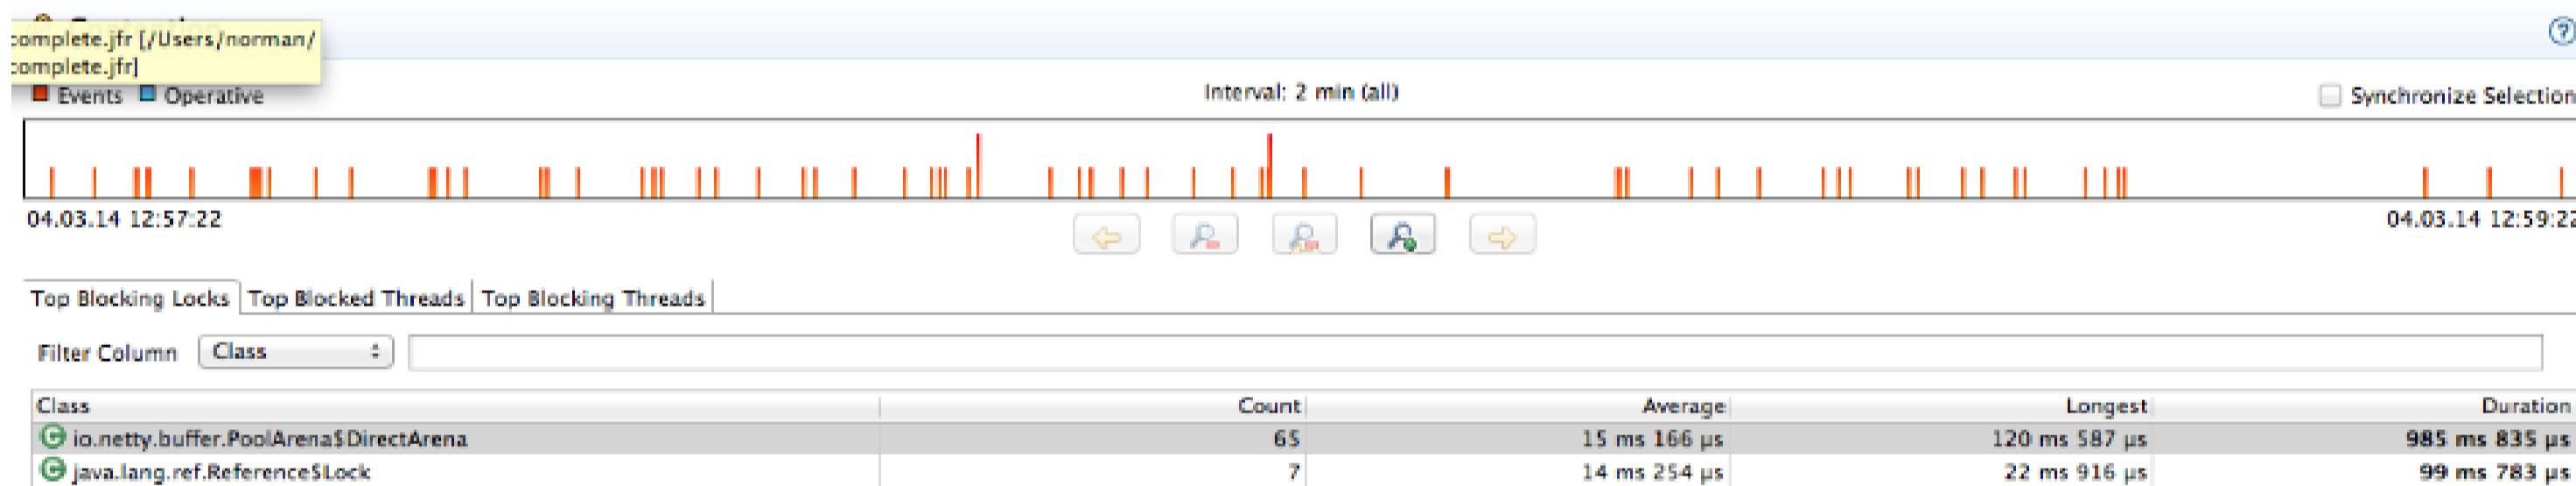
```
[nmaurer@xxx]~% wrk/wrk -H 'Host: localhost' -H 'Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8' -H 'Connection: keep-alive' -d 120 -c 256 -t 16 --pipeline 256 http://xxx:8080/plain/text  
...  
Requests/sec: 2812559.99  
Transfer/sec: 388.93MB
```



# PooledByteBufAllocator - caches to the rescue!

## With caches

```
[nmaurer@xxx]~% wrk/wrk -H 'Host: localhost' -H 'Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8' -H 'Connection: keep-alive' -d 120 -c 256 -t 16 --pipeline 256 http://xxx:8080/plain/text  
...  
Requests/sec: 3022942.17  
Transfer/sec: 418.02MB
```



# Leak detection built-in

## simple leak reporting

```
LEAK: ByteBuf.release() was not called before it's garbage-collected....
```

## advanced leak reporting

```
LEAK: ByteBuf.release() was not called before it's garbage-collected.
```

```
Recent access records: 1
```

```
#1:
```

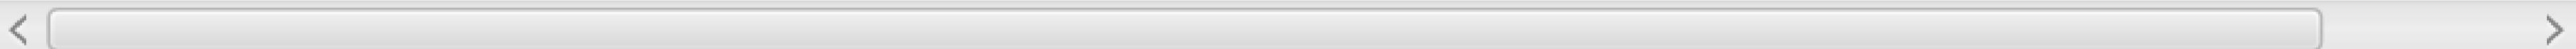
```
io.netty.buffer.AdvancedLeakAwareByteBuf.toString(AdvancedLeakAwareByteBuf.java:697)
```

```
...
```

```
Created at:
```

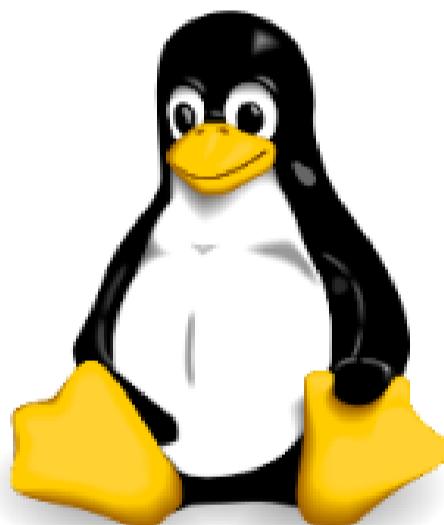
```
...
```

```
io.netty.handler.codec.xml.XmlFrameDecoderTest.testDecodeWithXml(XmlFrameDecoderTest.ja
```



## Native stuff in Netty 4

- 💡 OpenSSL based SslEngine to reduce memory usage and latency.
- 💡 Native transport for linux using **epoll ET** for more performance and less CPU usage.
- 💡 Native transport also supports **SO\_REUSEPORT** and **TCP\_CORK** :)



# Switching to native transport is easy

## Using NIO transport

```
Bootstrap bootstrap = new Bootstrap().group(new NioEventLoopGroup());  
bootstrap.channel(NioSocketChannel.class);
```

## Using native transport

```
Bootstrap bootstrap = new Bootstrap().group(new EpollEventLoopGroup());  
bootstrap.channel(EpollSocketChannel.class);
```

---

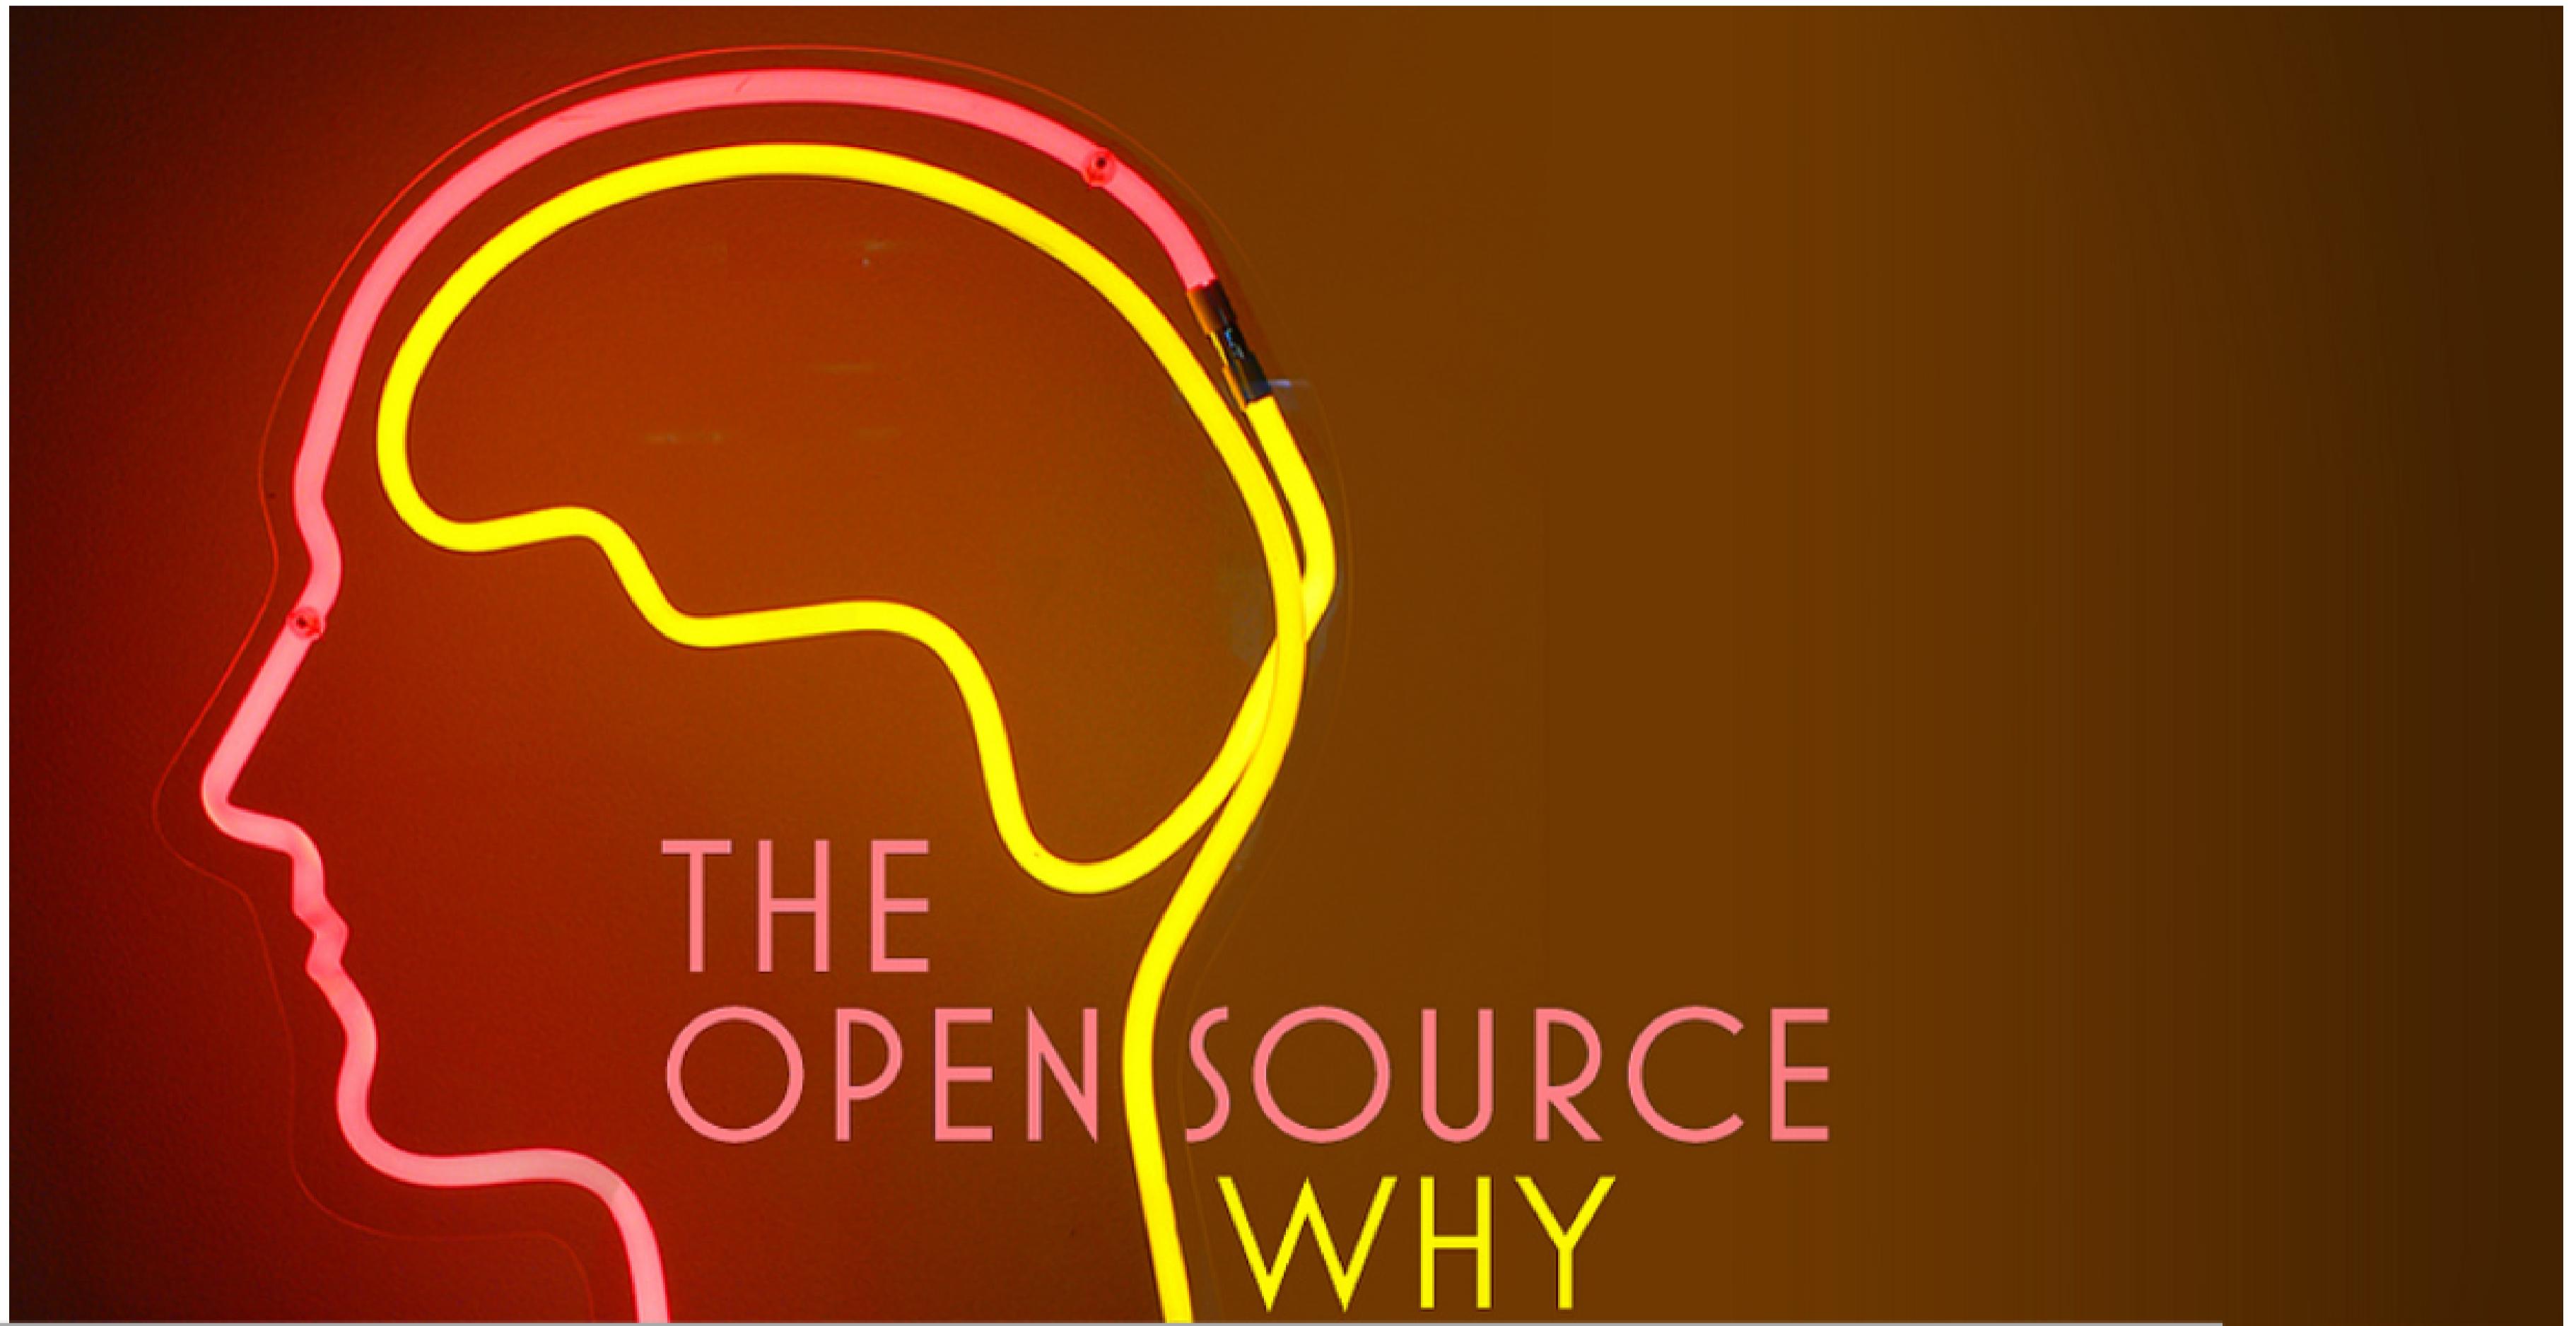
## Things to come

- Dynamic **Channel** re-register (based on metrics) - GSOC
- **ForkJoinPool** based **EventLoop** - GSOC
- Metrics



<https://www.flickr.com/photos/mrpinkeyes/5140672916/>

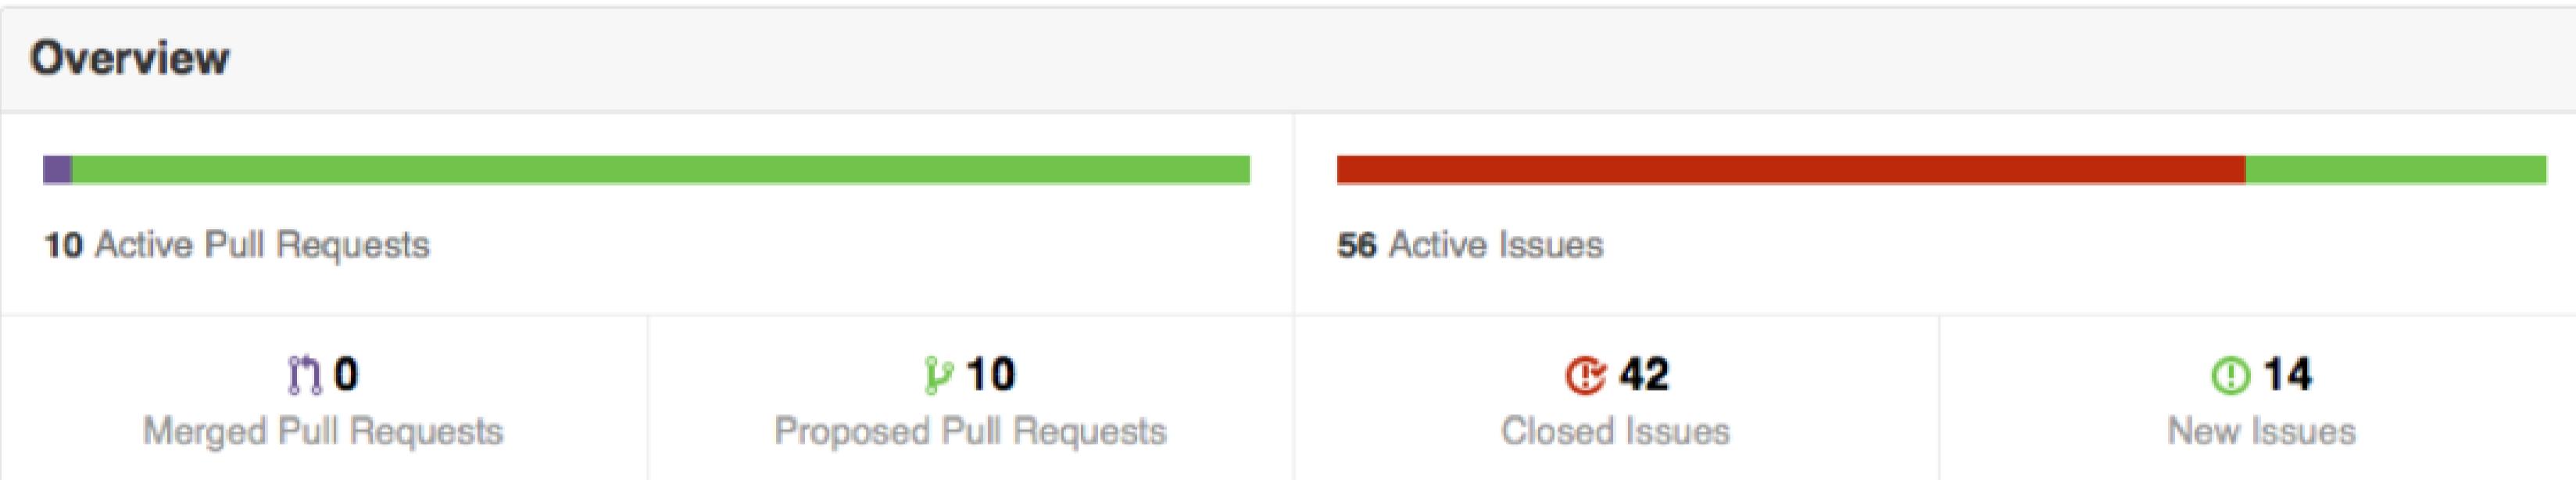
**It's OpenSource**



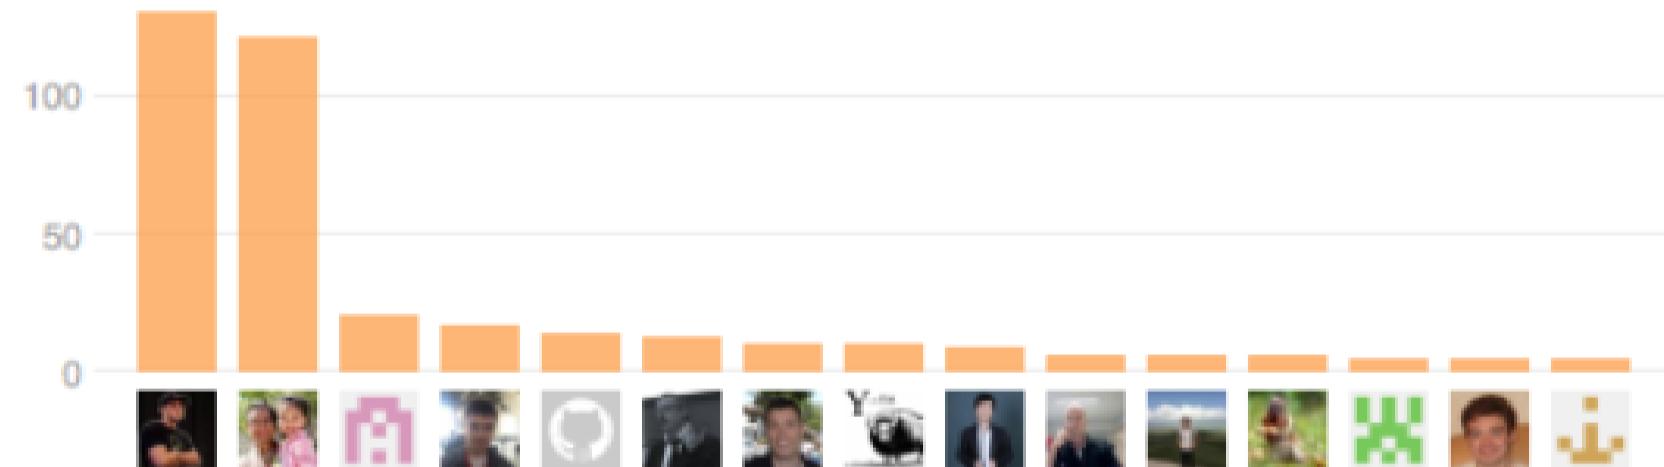
# Vibrant community!

June 08 2014 - July 08 2014

Period: 1 month ▾



Excluding merges, **25 authors** have pushed **102 commits** to master and **398 commits** to all branches. On master, **485 files** have changed and there have been **21,210 additions** and **6,740 deletions**.



## Companies using Netty!



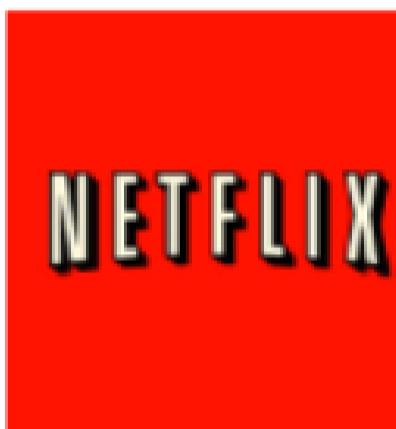
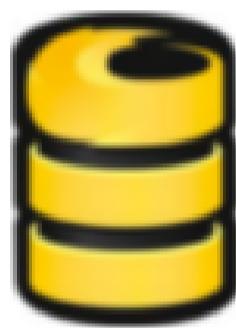
redhat.



cisco



ARISTA



be free  
boundary

Google



KakaoTalk



Typesafe



Couchbase

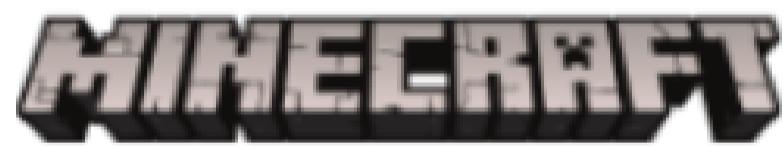


basho



... and many more ...

# (Opensource) Projects using Netty!



 ... and many more ...

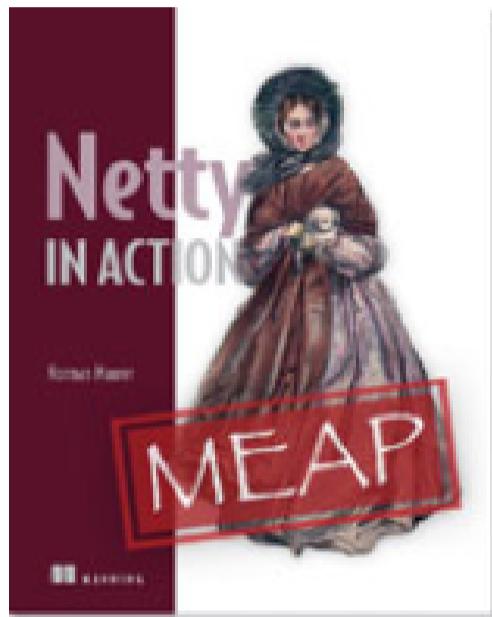
## Get Involved - We love contributions



- i** Mailinglist - <https://groups.google.com/forum/#!forum/netty>
  - i** IRC - #netty irc.freenode.org
  - i** Website - <http://netty.io>
  - i** Source / issue tracker - <https://github.com/netty/netty/>
-

**Want to know more?**

💡 Buy my book [Netty in Action](#) and make me **RICH**.



<http://www.manning.com/maurer>

**\$ KA-CHING \$**

---

## References

- Netty - <http://netty.io>
  - Slides generated with Asciidoctor and DZSlides backend
  - Original slide template - Dan Allen & Sarah White
  - All pictures licensed with Creative Commons Attribution or Creative Commons Attribution-Share Alike
-

# Norman Maurer



@normanmaurer